



Lab1



default ▾

%spark

FINISHED ▶ ⌵ ⌵ ⌵ ⌵

```
import org.apache.spark.rdd._
import scala.collection.JavaConverters._
import au.com.bytecode.opencsv.CSVReader

import java.io._
import org.joda.time._
import org.joda.time.format._
import org.joda.time.format.DateTimeFormat
import org.joda.time.DateTime
import org.joda.time.Days
```

```
import org.apache.spark.rdd._
import scala.collection.JavaConverters._
import au.com.bytecode.opencsv.CSVReader
import java.io._
import org.joda.time._
import org.joda.time.format._
import org.joda.time.format.DateTimeFormat
import org.joda.time.DateTime
import org.joda.time.Days
```



Took 4 sec. Last updated by anonymous at February 03 2017, 1:56:49 PM.

```
case class DelayRec(year: String,
                    month: String,
                    dayOfMonth: String,
                    dayOfWeek: String,
                    crsDepTime: String,
                    depDelay: String,
                    origin: String,
                    distance: String,
                    cancelled: String) {
```

FINISHED ▶ ⌵ ⌵ ⌵ ⌵

```
val holidays = List("01/01/2007", "01/15/2007", "02/19/2007", "05/28/2007", "06/07/2007",
                    "09/03/2007", "10/08/2007", "11/11/2007", "11/22/2007", "12/25/2007",
                    "01/01/2008", "01/21/2008", "02/18/2008", "05/22/2008", "05/26/2008", "07/04/2008",
                    "09/01/2008", "10/13/2008", "11/11/2008", "11/27/2008", "12/25/2008")
```

```
def gen_features: (String, Array[Double]) = {
  val values = Array(
    depDelay.toDouble,
    month.toDouble,
    dayOfMonth.toDouble,
    dayOfWeek.toDouble,
    get_hour(crsDepTime).toDouble,
    distance.toDouble,
    days_from_nearest_holiday(year.toInt, month.toInt, dayOfMonth.toInt)
  )
  new Tuple2(to_date(year.toInt, month.toInt, dayOfMonth.toInt), values)
```

Lab 1

Untitled Untitled Untitled Untitled Untitled Untitled Untitled Untitled

Zeppelin

```
def get_hour(deptime: String) : String = "%04d".format(deptime.toInt).take(2)
def to_date(year: Int, month: Int, day: Int) = "%04d%02d%02d".format(year, month, day)
```

Lab 1

```
def days_from_nearest_holiday(year: Int, month: Int, day: Int): Int = {
  val sampleDate = new org.joda.time.DateTime(year, month, day, 0, 0)

  holidays.foldLeft(3000) { (r, c) =>
    val holiday = org.joda.time.DateTimeFormat.forPattern("MM/dd/yyyy").parseDateT
    val distance = Math.abs(org.joda.time.Days.daysBetween(holiday, sampleDate).getDays)
    math.min(r, distance)
  }
}
```

default

```
warning: Class org.joda.convert.FromString not found - continuing with a stub.
warning: Class org.joda.convert.ToString not found - continuing with a stub.
warning: Class org.joda.convert.ToString not found - continuing with a stub.
warning: Class org.joda.convert.FromString not found - continuing with a stub.
warning: Class org.joda.convert.ToString not found - continuing with a stub.
warning: Class org.joda.convert.FromString not found - continuing with a stub.
warning: Class org.joda.convert.ToString not found - continuing with a stub.
defined class DelayRec
```

Took 4 sec. Last updated by anonymous at February 03 2017, 1:57:02 PM.



%sql

ERROR

```
select dayofWeek, case when depDelay > 15 then 'delayed' else 'ok' end , count(1)
from data_2007tmp
group by dayofweek , case when depDelay > 15 then 'delayed' else 'ok' end
```

```
Table or view not found: data_2007tmp; line 2 pos 5
set zeppelin.spark.sql.stacktrace = true to see full stacktrace
```

Took 13 sec. Last updated by anonymous at February 03 2017, 1:32:03 AM.



%sh

FINISHED

```
wget http://stat-computing.org/dataexpo/2009/2008.csv.bz2 -O /Users/jyothi/Downloads/flights_
```

```
--2017-02-03 14:04:36-- http://stat-computing.org/dataexpo/2009/2008.csv.bz2
Resolving stat-computing.org... 52.218.160.39
Connecting to stat-computing.org|52.218.160.39|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 113753229 (108M) [application/x-bzip2]
Saving to: '/Users/jyothi/Downloads/flights_2008.csv.bz2'
```



```
0K ..... 0% 322K 5m45s
50K ..... 0% 625K 4m21s
100K ..... 0% 627K 3m53s
150K ..... 0% 678K 3m36s
200K ..... 0% 5.55M 2m56s
250K ..... 0% 818K 2m49s
300K ..... 0% 2.33M 2m32s
350K ..... 0% 9.04M 2m14s
400K ..... 0% 12.2M 2m0s
450K ..... 0% 885K 2m1s
500K ..... 0% 4.34M 1m52s
```

Lab 1

Untitled Untitled Untitled Untitled Untitled Untitled Untitled Untitled Untitled Untitled

Zeppelin

Lab 1



ERROR ▶ ⌵ ⌵ ⌵ ⌵ ⌵
default ▼

```
#remove existing copies of dataset from HDFS
hadoop fs -rm -r -f /tmp/airflightsdelays
```

```
17/02/03 13:48:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
rm: Call From Jyothi.local/172.17.204.216 to localhost:9000 failed on connection exception: java.net.ConnectException: Connection refused; For more details see: http://wiki.apache.org/hadoop/ConnectionRefused
17/02/03 13:48:34 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
rm: Call From Jyothi.local/172.17.204.216 to localhost:9000 failed on connection exception: java.net.ConnectException: Connection refused; For more details see: http://wiki.apache.org/hadoop/ConnectionRefused
ExitValue: 1
```

Took 8 sec. Last updated by anonymous at February 03 2017, 1:48:37 PM.

FINISHED ▶ ⌵ ⌵ ⌵ ⌵ ⌵

```
// function to do a preprocessing step for a given file
def prepFlightDelays(infile: String): RDD[DelayRec] = {
  val data = sc.textFile(infile)

  data.map { line =>
    val reader = new au.com.bytecode.opencsv.CSVReader(new StringReader(line))
    reader.readAll().asScala.toList.map(rec => DelayRec(rec(0),rec(1),rec(2),rec(3),rec(5),rec(6)))
  }.map(list => list(0))
  .filter(rec => rec.year != "Year")
  .filter(rec => rec.cancelled == "0")
  .filter(rec => rec.origin == "ORD")
}

val data_2007tmp = prepFlightDelays("/Users/jyothi/Downloads/flights_2007.csv")
val data_2007 = data_2007tmp.map(rec => rec.gen_features._2)
val data_2008 = prepFlightDelays("/Users/jyothi/Downloads/airflightsdelays/flights_2008.csv")

data_2007tmp.toDF().registerTempTable("data_2007tmp")

data_2007.take(5).map(x => x.mkString ",").foreach(println)
```

```
prepFlightDelays: (infile: String)org.apache.spark.rdd.RDD[DelayRec]
data_2007tmp: org.apache.spark.rdd.RDD[DelayRec] = MapPartitionsRDD[64] at filter at <console>:58
data_2007: org.apache.spark.rdd.RDD[Array[Double]] = MapPartitionsRDD[65] at map at <console>:52
data_2008: org.apache.spark.rdd.RDD[Array[Double]] = MapPartitionsRDD[73] at map at <console>:50
warning: there was one deprecation warning; re-run with -deprecation for details
-8.0,1.0,25.0,4.0,11.0,719.0,10.0
41.0,1.0,28.0,7.0,15.0,925.0,13.0
45.0,1.0,29.0,1.0,20.0,316.0,14.0
-9.0,1.0,17.0,3.0,19.0,719.0,2.0
180.0,1.0,12.0,5.0,17.0,316.0,3.0
```

Took 5 sec. Last updated by anonymous at February 03 2017, 2:07:09 PM.

Lab 1

```

yearofweek, case when depDelay > 15 then 'delayed' else 'ok' end, count(1)
2007tmp
dayofweek, case when depDelay > 15 then 'delayed' else 'ok' end

```

default ▼



dayofWeek	CASE WHEN (CAST(depDelay AS DOUBLE) > CAST(15 AS DOUBLE)) THEN delayed ELSE ok
1	delayed
7	ok
1	ok
6	delayed
2	delayed
3	ok
4	delayed
3	delayed

Took 26 sec. Last updated by anonymous at February 03 2017, 2:07:43 PM.

```
%sql
```


FINISHED    

```
select cast( cast(crsDepTime as int) / 100 as int) as hour, case when depDelay > 15 then 'delayed'
count from data_2007tmp
group by cast( cast(crsDepTime as int) / 100 as int), case when depDelay > 15 then 'delayed'
```



hour	delay	count
12	ok	13
13	ok	20
20	delayed	10
10	ok	17
19	ok	12
15	ok	14
15	delayed	7,1
21	ok	8,6

Took 27 sec. Last updated by anonymous at February 03 2017, 2:09:43 PM.



Zeppelin

READY ▶ ⌵ ⌶ ⚙

Lab1

▶ ⌵ ⌶ ✎ 📄 ⬇ 📄

🗑 ⏻

⌨ ⚙ 🔒 default ▼